

# Logistic Regression

One sample,  $X \Rightarrow x_1, x_2$  (2 features)

which class do these features correspond to??

$\Rightarrow$  0/1??

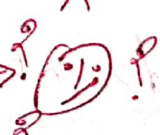
$\hookrightarrow$  binary classification

$$(1, x_1, x_2) \cdot \text{dot}(\theta_3, \theta_1, \theta_2) \Rightarrow \textcircled{X\theta}$$
$$= \theta_3 + \theta_1 x_1 + \theta_2 x_2$$

$\hookrightarrow$  we update and tune  $\theta_1, \theta_2$  and  $\theta_3$  such that  $X\theta$  becomes close to 0 for  $x_1, x_2$  coming from class 0 samples, close to 1 for  $x_1, x_2$  coming from class 1 samples.

$\theta \Rightarrow (\theta_3, \theta_1, \theta_2) \Rightarrow$  parameters  
           $\downarrow$            $\downarrow$   
          bias     weights

Its difficult to get  $X\theta$  in the range of  $[0, 1]$ .

Activation function  $\rightarrow$    $g(z)$

$$z = X\theta = \theta_3 + \theta_1 x_1 + \theta_2 x_2$$

$$h = g(z) = \frac{1}{1 + e^{-z}} \text{ (sigmoid function)}$$

$h \Rightarrow$  predicted value  $\Rightarrow [0, 1]$

$y \Rightarrow$  true label

Now how do we know what values to use for  $\theta$ ??

Ans: by gradient descent optimization

$$J(\theta) = -y \log(h) - (1-y) \log(1-h)$$

$\rightarrow$  log loss function

$$y=0, h=0 (\text{correct}) \Rightarrow -0 \cdot \log(0) - 1 \cdot \log(1) \rightarrow 0$$

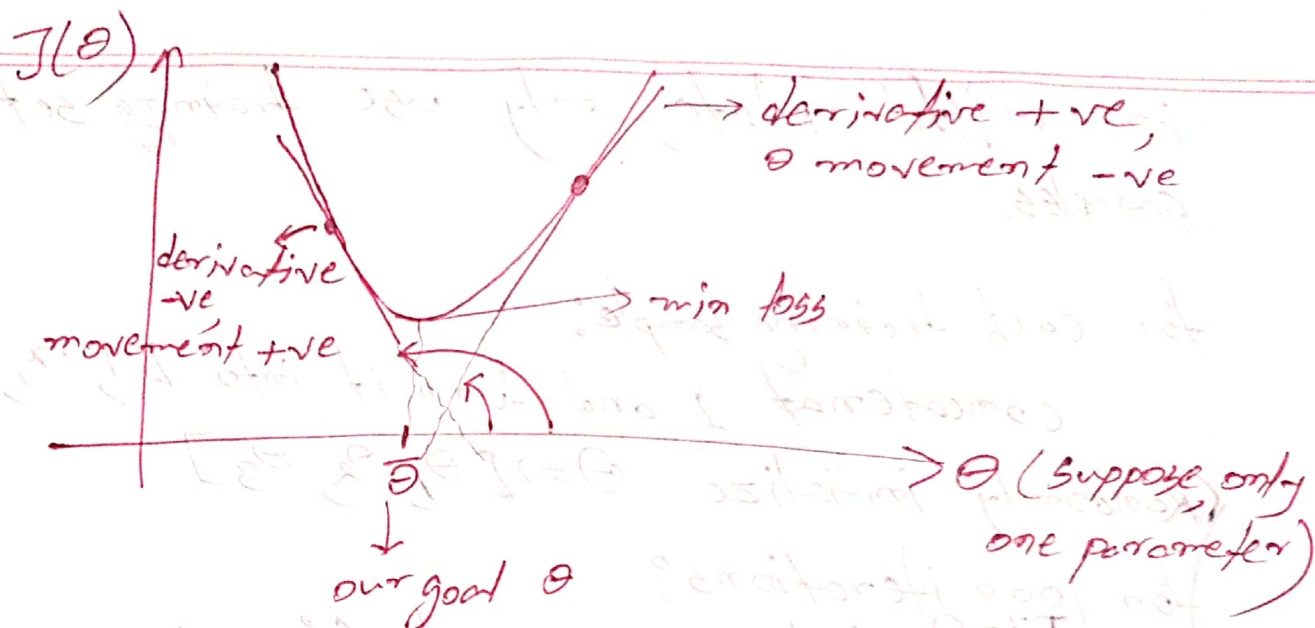
$$y=1, h=1 (\text{correct}) \Rightarrow -1 \cdot \log(1) - 0 \cdot \log(0) \rightarrow 0$$

$$y=1, h=0 (\text{wrong}) \Rightarrow -1 \cdot \log(0) - 0 \cdot \log(1) \Rightarrow -(-\infty) - 0 \Rightarrow +\infty$$

The closer the value of  $h$  to  $y$ , less will be the loss.

⊙!! how do we update  $\theta$ ??





$$dv = \frac{\partial J(\theta)}{\partial \theta} = x(h - y) \Rightarrow \text{calculate and see for yourself}$$

$$\theta = \theta - dv$$

↳ a slight problem.

What if  $dv$  is too BIG??  
will never reach  $\bar{\theta}$ .

learning rate (l.r.) is the answer

$$\theta = \theta - dv * \text{l.r.} \rightarrow \text{l.r.} < 1 \text{ and } \text{l.r.} > 0$$

typically  $\Rightarrow 0.01, 0.001, 0.0001$   
we set this manually  
by looking at validation results  
for different l.r.

How many times to update  $\theta$ ?  $\hookrightarrow$  hyperparameter  
↳ until loss function value does not decrease any longer.

For weight update, only use training set samples.

for each training sample:

concatenate 1 and turn it into  $[1, x_1, x_2]$

Randomly initialize  $\Theta \Rightarrow [\theta_1, \theta_2, \theta_3]$

for 1000 iterations:

$TJ = 0$   
for each training sample  $s$ :

$z = x \cdot \Theta$  [ $x \Rightarrow s$  features,  $\Theta \Rightarrow$  parameters]  
use np.dot]

$h = \text{sigmoid}(z)$   $\rightarrow$  available in python library

$$J = -y \log(h) - (1-y) \log(1-h)$$

$\rightarrow$  true label of  $s$

$$TJ = TJ + J$$

$\rightarrow$  total loss for this iteration

$$dv = x(h - y)$$

$\rightarrow$  it has 3 dimensions

$$\Theta = \Theta - dv * lr$$

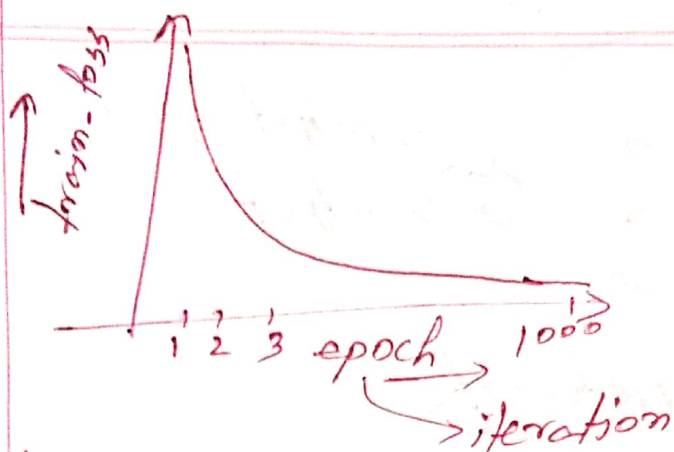
$\rightarrow$  also has 3 dimensions

$$TJ = TJ / (\text{training sample no.})$$

Append  $TJ$  in a list named "train-loss"

training  $\Rightarrow$  updating  $\Theta$





validation

correct = 0  
for each validation sample  $s^o$ :

$$z = X \theta$$

$$h = \text{sigmoid}(z)$$

if  $h \geq 0.5$ :

$$\text{~~states~~ } h = 1$$

else:  $h = 0$

if  $h = y$ :

$$\text{correct} += 1$$

$$\text{val-acc} = \text{correct} / (\text{total val sample no.}) \times 100$$

perform this training-validation for  $\text{l.r.} = 0.1, 0.01, 0.001, 0.0001, 0.00001$  and report val-acc in a table. → hyperparameter tuning

Take the  $\text{l.r.}$  with best val-acc

Use this  $\text{l.r.}$  for training and determine

test-acc with the trained model

